

Conor Reisman

4/12/2013

CSE444

Lab1

## Design Decisions

Tuple and TupleDesc were both implemented as lists. This makes the class work in  $O(1)$  time for every method other than fieldNameToIndex (not including the ones that must be  $O(N)$ ). I could make it  $O(1)$  with a map, but it's not worth it.

Catalog is a pair of maps, one from names to ids and one from ids to relevant data (DbFile, primary key, and name). This is  $O(1)$  for all queries.

BufferPool is a map from PageId to Page.

HeapPage has a map from Tuple to its index in the Tuple array. It also has a list of empty slots for inserting new Tuple.

HeapFile has a static counter to assign unique ids. These are cached by file name in a static map so HeapFiles with the same name won't get different ids. As such, I didn't follow the recommended implementation of hashing.

## API Changes

Most API changes were very minor, such as throwing IllegalArgumentExceptions on invalid inputs. Other changes made were adding implements Iterator/Iterable to different classes (Tuple, TupleDesc, and HeapPage). The most important change was to make the Page interface implement Iterable. This allows me to not have to cast Page to HeapPage in order to get an iterator in HeapFile. It also makes sense since all pages should have an iterator method that returns an Iterator<Tuple>.

## Incomplete/Missing Implementation

Nothing that was required for lab 1 was left unimplemented.

## Other

I spent around 15 hours implementing and reviewing this assignment.

A list of my complaints about this assignment can be found in notes.txt.